

The details of the contribution of IJDL special issue submission #XXXX

Not All Mementos Are Created Equal: Measuring The Impact Of Missing Resources

Section	DL2014 Proceeding	IJDL Special Issue
Abstract		Added results of the additional experiments and investigations
Conclusion		Added outcomes and impacts of the additional experiments and investigations
Dataset		
Archives	Internet Archive	To complement our discussion of the damage in the Internet Archive, we measured a page-at-a-time archiver, WebCite. We also included a discussion of archive.today and compared the three archives.
Mementos	45,341 URI-Ms <i>from IA</i>	WebCite had 992 mementos to add to the measurements, as we submitted the URI-Rs in our dataset to Archive.today and Webcite and created two additional sets of 1,861 mementos to evaluate.
Measurements		
Damage over time	Measurement and discussion of damage in the Internet Archive	Comparison of memento damage and missing resources over time for WebCite, and compared the archive performance to that of the Internet Archive
Collection Damage	Measurement and discussion of damage in the Internet Archive	Discussion of the impact of missing embedded resources on the overall damage of WebCite
Number of missing resources	Measurement and discussion of missing embedded resources in the Internet Archive	Discussion of the number of missing embedded resources, their impact on damage, a comparison to the Internet Archive, and a discussion of what we can deduce from missing resources and damage
Missing File types	Measurement and discussion of missing file types in the Internet Archive	Comparison of the differences in archiving successes in the Internet Archive and WebCite, as well as the impact this may have on future damage measurements

$$\begin{aligned} DA &= 45,341 \\ Webcite &= 992 \\ AT &= \underline{\quad} \end{aligned}$$

Additional Discussion			
Impact of JavaScript	No discussion of JavaScript and the impact on damage	A discussion of how the ability or inability to execute JavaScript during archiving can impact the damage of mementos	
Page-at-a-time archivers	No discussion of page-at-a-time archivers	Discussion of page-at-a-time archivers (WebCite and Archive.today) and how they differ from traditional crawler-based archivers	
PhantomJS	No discussion of JavaScript and the impact on Damage	Included an experiment to measure the impact that JavaScript has on damage, and how Archive.today provides an improvement to WebCite, presumably due to its ability to handle JavaScript	
Document Statistics			
Words	7,275	9,909	+36%
Figures	7	14	+100%
Sections & subsections	11	16	+45%
Evaluations	Archives: 1 Mementos: 45,341	Archives: 3 Mementos: +992 + $2 * 1,861 = 4,714$ new mementos $45,341 + 4,714 =$ 50,055	+200% +10%
Pages	10	17	+70%
Tables	10	10	+0%
References	28	36	+29%

Not All Mementos Are Created Equal: Measuring The Impact Of Missing Resources

Justin F. Brunelle, Mat Kelly, Hany SalahEldeen, Michele C. Weigle, and Michael L. Nelson

Received: date / Accepted: date

Abstract Web archives do not always capture every resource on every page that they attempt to archive. This results in archived pages missing a portion of their embedded resources. These embedded resources have varying historic, utility, and importance values. The proportion of missing embedded resources does not provide an accurate measure of their impact on the Web page; some embedded resources are more important to the utility of a page than others. We propose a method to measure the relative value of embedded resources and assign a damage rating to archived pages as a way to evaluate archival success. In this paper, we show that Web users' perceptions of damage are not accurately estimated by the proportion of missing embedded resources. In fact, the proportion of missing embedded resources is a less accurate estimate of resource damage than a random selection. We propose a damage rating algorithm that provides closer alignment to Web user perception, providing an overall improved agreement with users on memento damage by 17% and an improvement by 51% if the mementos have a damage rating delta > 0.30. We use our algorithm to measure damage in the Internet Archive, showing that it is getting better at mitigating damage over time (going from a damage rating of 0.16 in 1998 to 0.13 in 2013). However, we show that a greater number of important embedded resources (2.05 per memento on average) are missing over time. Alternatively, the damage in WebCite is increasing over time (going from 0.375 in 2007 to 0.475 in 2014) while the missing embedded resources is remaining constant (13% of the resources are miss-

ing on average). Finally, we investigate the impact of JavaScript on the damage of the archives, showing that a crawler that can archive JavaScript dependent representations will improve memento damage by 13.5%.

Keywords Web Architecture, Web Archiving, Digital Preservation

1 Introduction

Web archives are valuable cultural repositories that capture and store Web content. ~~Users make use of~~ archives like the Internet Archive [21, 33] to retrieve archived material [14, 19] for a variety of purposes and in a variety of ways [2]. However, the resources being requested by Web users may not be complete; embedded resources are sometimes missing from an archived Web page [4]. Missing embedded resources return a non-200 HTTP status (e.g., 404, 503) when their URI is dereferenced.

Large images are often more important to an archived page's utility than small images. Similarly, stylesheets that format visible content are more important to the representation of the page than stylesheets without significant formatting responsibilities. We provide a mechanism to assess the impact of missing embedded resources in the archives.

Throughout this paper we use Memento Framework terminology. Memento [34] is a framework that allows ~~web users to browse in the temporal dimension by aggregating the offerings of the archives at a single point of access~~. Original (or live web) resources are identified by URI-R, and archived versions of URI-Rs are called *mementos* and are identified by URI-M. Memento TimeMaps are machine-readable lists of mementos (at the level of single-archives or aggregation-of-archives) sorted by archival date.

Old Dominion University, Department of Computer Science
Norfolk VA, 23529, USA
{jbrunelle,mkelly,hany,mweigle,mln}@cs.odu.edu

People (and robots)
standardizes
web archive
access
and
terminology

This research is motivated by three factors. First, we want to understand how missing embedded resources impact Web user satisfaction (i.e., the utility of mementos). Using an algorithm to measure embedded resource importance, we determine whether an important embedded resource of the memento is missing (e.g., a main image or video essential to the user's understanding of the page), or the missing embedded resource is a spacer image or a small button logo that contributes little to the memento's utility for the user. We propose a method of weighting embedded resources in a memento according to importance. We show that this is an improved damage rating over an unweighted count of missing embedded resources. We use Amazon's Mechanical Turk to compare our algorithm to Web users' notion of damage and to show an improvement over the unweighted count of missing embedded resources.

Second, we use our algorithm to assess the damage of mementos in the Internet Archive. We use the unweighted measure of damage as the proportion of missing embedded resources to all requested resources (M_m) and compare it to our calculation of damage (D_m).

Third and finally, we measure damage in the Internet Archive over time using our weighted algorithm. We then describe how this algorithm can be used for future enhancements of the Heritrix crawler [20, 28] and archival processes.

2 Motivating Examples

We use the XKCD Web page as an example of a resource with embedded resources of differing importance. We captured the URI-R using the wget [12] command¹ and manually inflicted damage on a local memento of <http://www.xkcd.com/> by removing embedded images. We used PhantomJS [23] to dereference the URI-M, take a PNG snapshot of the representation, and recorded the resulting HTTP response headers of the embedded resources. We created three mementos of the URI-R: one duplicating its live Web counterpart (m_0), one with the central comic image removed (m_1), and one with two logo images removed (m_2). The snapshots taken by PhantomJS are provided in Figures 1(a), 1(b), and 1(c). As shown in the captions, the proportion of embedded missing resources (M_m) varies among the mementos.

At the time of this test, the live XKCD site was missing two embedded stylesheets, as are m_0 , m_1 , and m_2 since they are copies of the live site. We verified that our memento m_0 has a M_m value identical to its live Web counterpart – the live resource and m_0 are both

missing the same embedded resources ($M_m=0.17$). In Figure 1(a), m_0 has multiple embedded resources, but we focus on the three identified by the red arrows: the XKCD logo, the main comic image, and the banner of comics. The central image is most important to the utility of the page – without the main comic image, the user does not obtain the information from the page that the author intended (Figure 1(b)). The logo and banner are not essential to the user's understanding of the XKCD content (Figure 1(c)).

Cascading Stylesheets (CSS) also differ in importance. Some stylesheets are responsible for formatting small portions of a page, while others are responsible for placing images and other content or even organizing the entire page for the user. Figure 1(d) shows a memento of a URI-R that is missing a single stylesheet. This stylesheet is responsible for a large amount of information in the representation and without it, the meaning and utility of the memento changes. Figure 1(e) shows a memento that is properly styled but is missing two stylesheets that are not responsible for the majority of the content organization.

As we have discussed, the percentage of successfully dereferenced embedded resources is not the only factor in determining memento quality. In support of that principle, we refer to Figure 1(e) in which $M_m=0.2$ (6/30). However, it appears to be well-preserved. In our XKCD example, Figure 1(c) is missing two images ($M_m=0.24$) yet maintains more important embedded mementos than Figure 1(b) ($M_m=0.29$). These examples support the motivation of our research by demonstrating that unweighted percentages (i.e., M_m) are insufficient to assess perceived memento damage.


3 Related Work

Past researchers have studied the completeness of the archives, the recrawl policies that optimize archive quality, and the relative importance of information of content within Web resources. We build upon these prior works and apply their findings to develop our algorithm for automatically assess the quality of mementos.

SalahEldeen et al. have studied the rate at which live resources disappear from the Web. In a study of the Egyptian Revolution, SalahEldeen found that 11% of the resources shared over Twitter were missing after one year [26, 27].

Our previous work studied the factors influencing archivability, including accessibility standards and their impact on memento completeness [16]. In this work, Kelly used a yearly sampling method to select mementos for testing. We use a similar method in this work to study memento damage.


¹ We executed the wget command with parameters as follows: wget -E -H -K -p <http://www.xkcd.com/>



(a) All three of the embedded images are included in m_0 and identified by the red arrows ($M_m=0.17$).

(b) We removed the large, central image (that is the main content of the page) from m_1 , identified by the red arrow ($M_m=0.24$).

(c) We removed the XKCD logo and banner of comics from m_2 , identified by the red arrows ($M_m=0.29$).



(d) This memento (URI-M <http://web.archive.org/web/20110116022653/http://www.cityofmoorhead.com/flood/?>) is missing two stylesheets which changes the entire appearance and utility of the memento ($M_m=0.38$).

(e) Meanwhile, this memento (URI-M <http://web.archive.org/web/2006102083228/http://www.ascc.edu/>) is missing two stylesheets (along with two images) but does not appear damaged ($M_m=0.20$).

Fig. 1 Mementos have different meanings and usefulness depending on which embedded resources are missing from the memento (and the proportion of missing resources, M_m).

Spaniol et al. has measured the quality of Web archives based on matching crawler strategies with resource change rates [8, 31, 32]. Ben Saad and Gançarski performed a similar study regarding the importance of changes on a page [5]. Gray and Martin created a framework for high quality mementos and assessing their quality by measuring the missing embedded resources [13]. While these studies focused on memento completeness and site coverage, we focus on assessing the importance of the artifacts that are missing.

Fersini et al. studied the importance of information blocks of a rendered Web page, finding that blocks with more images are more important [11]. Singh et al. found that multimedia within a page is essential for user understanding [29]. Ye et al. found that the information blocks close to the center of the viewport contain important information, while “noise” – or unimportant content – occurs on the fringes or edges of the page [35]. Kohlschütter et al. also found that important content was located in the center of pages [18]. Centrality is a way for authors to convey importance of informa-

related info → add their
or crawling strategies. IJDL 2012 (save as TCOL 2011)
+ TCOL 2011

tion to their users. For example, images in the center of the viewport are more important or contribute to the users' understanding of a page than those positions on the fringes or outside the viewport of a page. Using these prior findings, we constructed an algorithm to assess the importance of embedded resources based on their MIME type, location in the viewport, and size in pixels.

Banos et al. created an algorithm to evaluate archival success based on adherence to standards for the purpose of assigning a resource archivability score [4]. Zhang et al. studied human perception and human ability to recognize differences in images effectively determining human perception limitations for images at the pixel level [36]. Rademacher et al. used human perception to identify the visual factors that distinguished computer generated images from photographs [24]. We use human perception in a similar way to identify levels of memento damage.

The algorithm proposed in this paper determines the importance of embedded resources. Song et al. outlined an algorithm for determining the importance of sections of Web pages based on their content, size, and position [30]. We extend this algorithm (using many of the same principles) to measure the importance of missing embedded resources.

4 Users' Perception of Damage

As archivists, our perception of damage differs from that of more traditional Web users. To determine if M_m (percent missing) is a good estimate of human perception of damage, we used Amazon's Mechanical Turk to measure human agreement with M_m .

To ensure that Mechanical Turk workers (or more colloquially, "turkers") could evaluate damage, we presented turkers with pairs of mementos that had varying levels of damage and asked them to select the memento they preferred to keep if given a choice between the two.

We captured 11 hand-selected URI-Rs (Table 1) on a local server and created five versions of the mementos for each URI-R. We manually ~~inflicted~~ damage to the mementos to create the five categories of damage. For the category *missing image*, we removed a prominent image (empirically identified as important) from the memento. For the category *missing css*, we removed a prominent CSS file to cause formatting issues in the memento; we empirically selected the CSS file to remove based on the greatest human-perceived detrimental impact to the page layout. We also created the categories *missing all images* (we removed every embedded image), *missing all resources* (we removed all embedded

resources), and *original* (the URI-M was a direct copy of the live resource) and measured the M_m of each URI-M in each category. We refer to the four categories of damaged mementos in aggregate as m_1 and the *original* as m_0 . These categories created a variety of damage ratings by a variety of missing embedded resources for identical URI-Rs at an identical time point to provide a wide spectrum of mementos to be evaluated by turkers.



Fig. 2 We asked the turkers to select the less damaged of two mementos. *The two versions of the page are accessible in separate tabs.*

With the goal of determining whether or not turkers can recognize damage in a memento, we presented the turkers with a m_1 and its m_0 counterpart (that is, a "damaged" and its *ground-truth* memento) and asked the turkers "We saved two pages for you. For which page did we do a better job?" (Figure 2). For each URI-R, a pair of mementos consisting of m_0 and one of the

URI-R	m_0	M_m			
		missing image	missing css	missing all images	missing all
http://www.cs.odu.edu/~mln/	0.14	0.43	0.29	0.43	0.43
http://activehistory.ca/2013/06/myspace-is-cool-again-too-bad-they-destroyed-history-along-the-way/comment-page-1/	0.0	0.32	0.32	0.57	0.85
http://www.albop.com/	0.0	0.13	0.0	0.50	0.50
http://www.cs.odu.edu/	0.10	0.13	0.11	0.82	0.81
http://ws-dl.blogspot.com/2013/08/2013-07-26-web-archiving-and-digital.html	0.07	0.08	0.08	0.13	0.14
http://www.cnn.com/2013/08/19/tech/social-media/zuckerberg-facebook-hack/	0.19	0.22	0.28	0.46	0.57
http://xkcd.com/	0.14	0.38	0.31	0.53	0.54
http://www.mozilla.org/	0.80	0.80	0.80	0.877	0.89
http://www.eshow.com/	0.05	0.05	0.06	0.11	0.33
http://google.com/	0.0	0.0	0.0	0.0	1.0
http://php.net/	0.32	0.33	0.33	0.37	0.37

Table 1 The 11 URI-Rs used to create the manually damaged dataset. M_m values are provided for each m_1 .

four categories of m_1 were evaluated by five turkers for a total of 280 evaluations.

We show the judgement splits from the turker evaluations in Table 2. The judgement splits refer to the number of turkers that selected the correct-incorrect version. For example, a 0-5 split means all five turkers selected the m_1 (an incorrect selection), a 5-0 split means all five turkers selected the m_0 memento (the correct selection), and a 3-2 split means three turkers selected the m_0 memento and two selected the m_1 (a correct selection by the majority, but still a split decision among the turkers). For the purposes of this paper, we consider only 5-0 and 4-1 splits as agreement and all other splits as disagreement. ΔM_m refers to the delta between M_{m_0} and M_{m_1} .

The turkers selected m_0 as the preferred option (less damaged memento) 81% of the time (226/280). As ΔM_m shrinks, turker agreement is more consistent.

Regardless of ΔM_m , 81% of the evaluations agreed with M_m as a suitable damage metric (5-0 and 4-1 splits). Turkers were unsure about the damage (3-2 and 2-3 splits) 18% of the time and incorrectly identified damage only once. The average ΔM_m for the unsure selections was < 0.01, and the only 0-5 split had a ΔM_m of 0.014, suggesting that confusion or disagreement occurs more often when the damage delta is smaller.

Confusion matrices provide a consolidated view of an algorithm's performance. The top left quadrant shows the number of true positives, the top right shows the number of false negatives, the bottom left shows false positives, and the bottom right shows true negatives. The algorithm's accuracy ((True Positives + True Negatives) / (All Positives and Negatives)) and harmonic

ΔM_m	Splits						Total
	5-0	4-1	3-2	2-3	1-4	0-5	
1.0							0.00
0.9							0.00
0.8	4						0.07
0.7							0.00
0.6							0.00
0.5	1	1					0.04
0.4							0.00
0.3	15	5					0.36
0.2	2						0.04
0.1	5	4	4	2		1	0.29
0.0	5	3	1	3			0.22
Total	0.58	0.23	0.09	0.09	0.00	0.02	1.0

Table 2 The turkers selected m_0 as the preferred memento 81% of the time, and more consistently for larger ΔM_m values.

Turker Assessment	M_m	
	Select m_0	Select m_1
m_0	44	0
m_1	11	0

Table 3 Confusion matrix of the turker assessments of the m_0 vs m_1 comparison test.

mean (or F_1 Score: $2 * \text{True Positives} / (2 * \text{True Positives} + \text{False Positives} + \text{False Negatives})$) are calculated using a confusion matrix. A harmonic mean provides an average (in this case, of the algorithm's success rate) and is sensitive to small values and outliers.

From the confusion matrix (Table 3), we can calculate $F_1=0.88$ and accuracy=0.80 for m_0 vs m_1 . Turker agreement does not match M_m 100% of the time with the m_0 vs m_1 test because of phenomena with aesthetics and human perception. Also, m_0 is often incomplete and, as a result, has $D_{m_0} > 0$ (see php.net in Table 3).

5 Evaluating Organic Damage

Turkers identified m_0 in the m_0 vs m_1 in a large majority (81%) of the comparisons. As a result, the turkers have shown that they can identify a damaged resource when presented a damaged and undamaged memento. Because they can identify damage in mementos, we used turkers to evaluate our measured damage of mementos found in the Internet Archive.

This experiment uses the same set of 2,000 URI-Rs as in our previous work [7], which was sampled from Twitter and Archive-It. The first dataset, the *Twitter* set, consists of Bitly URIs shared over Twitter. The second dataset, the *Archive-It* set, was sampled from Archive-It collections. The Archive-It collections are created and curated by human users often corresponding to a certain event (e.g., National September 11 Memorial Museum) or a specific set of Web sites (e.g., City of San Francisco). We discarded non-HTML representations (e.g., JPEG and PDF) from both sets for a final dataset of 1,861 URI-Rs. Non-HTML representations do not contribute to this study since they do not have embedded resources. There is no overlap between the two sets.

Using this set of URI-Rs, we measured the damage of one memento per year from the Internet Archive TimeMap of each of the 1,861 URI-Rs, resulting in 45,341 URI-Ms. We randomly selected a subset of 100 URI-Ms from this set. Similar to the evaluation in Section 4, we gave turkers two mementos (we will generalize these to m_2 and m_3) from consecutive years from the same TimeMap and asked the turkers to select the less damaged memento (“We saved two pages for you. For which page did we do a better job?”) as shown in Figure 2. Because m_2 and m_3 are observed from the Internet Archive, neither is considered a *ground-truth*. We measured the damage M_m of mementos in the Internet Archive and compared it to the turker perception of the utility of the mementos.

Contrary to the test in Section 4, as ΔM_m grows, the turkers are not as effective at selecting the less damaged memento (the splits are shown in Table 4). The turkers only agree with M_m 12% of the time and completely disagree with M_m (1-4 and 0-5 splits) 44% of the time. This discrepancy demonstrates that turker assessment of damage does not match M_m . Additionally, we see that the turkers performed well when comparing m_0 vs m_1 (original vs damaged) but struggle to compare m_2 vs m_3 (damaged vs damaged).

From the confusion matrix (Table 5), we can calculate the accuracy of turker selections of m_2 vs m_3 agreement with M_m is 0.46 with $F_1=0.55$. In a Receiver Operating Characteristic (ROC) curve [10], we calculated

ΔM_m	Splits						Total
	5-0	4-1	3-2	2-3	1-4	0-5	
1.0					1		0.01
0.9							0.00
0.8							0.00
0.7		1					0.01
0.6					1		0.01
0.5							0.00
0.4		1					0.01
0.3	1		3	4	1	2	0.11
0.2		5	6	5	12	9	0.37
0.1	4	5	10	11	15	3	0.48
0.0							0.00
Total	0.05	0.12	0.19	0.20	0.30	0.14	1.0

Table 4 The turker evaluations of the m_2 vs m_3 comparisons when using M_m as a damage measurement.

Turker Assesment	M_m	
	Select m_2	Select m_3
m_2	29	24
m_3	23	24

Table 5 Confusion matrix of the turker assessments of the m_2 vs m_3 comparison test against M_m .

Damage Calculation	AUC	F_1	Accuracy
M_m	0.472	0.55	0.46
M_{m_0}	0.789	0.88	0.80

Table 6 When compared to random, M_m performs worse than random selection and is worse than the optimal performance of m_0 vs m_1 .

the Area Under the ROC Curve (AUC) for the results of the turker evaluations of m_2 vs m_3 against M_m and the results of the manually damaged m_0 vs m_1 test (as the optimal performance). The AUC of M_m is lower (AUC=0.472) than random (AUC=0.500) as shown in Table 6, meaning that M_m performed worse than random for matching turker perception of damage and far worse than the optimal performance (AUC=0.789), a further indicator that M_m is not a suitable metric for measuring memento damage.

6 Calculating Memento Damage

With M_m not matching Web users’ perception of damage, we propose a new algorithm for assessing memento damage. Our proposed algorithm is based on the MIME type, size, and location of the embedded resource.

We define D_m as the damage rating, or cumulative damage, of a memento m in Equation 1. D_m is a normalized value ranging from [0, 1]. We calculate the potential damage of a memento and the actual damage of a memento and express the damage rating as the ratio of actual to potential damage. Notionally, potential damage is the cumulative importance of all embedded resources in the memento, while actual damage is only

↗ any natural
subsets
for
fn.3
section
one
big
section

the importance of those embedded resources that are unsuccessfully dereferenced, or missing.

$$D_m = \frac{D_{m_{actual}}}{D_{m_{potential}}} \quad (1)$$

To determine potential and actual damage, we first define the set of all embedded resources R and the set of all missing resources R_r in Equation 2.

$$\begin{aligned} R &= \{\text{All embedded resources requested}\} \\ R_r &= \{\text{All missing embedded resources}\} \\ R_r &\subseteq R \end{aligned} \quad (2)$$

We calculate the importance of each embedded resource in the set R . The sum of each embedded resource is the potential damage $D_{m_{potential}}$ (Equation 3). Important resources are assigned additional weights to increase their relative value over unimportant resources (Equations 5 - 6).

$$\begin{aligned} D_{m_{potential}} &= \frac{\sum_{i=1}^{n_{I,MM}} D_{[I,MM]}(i)}{n_{[I,MM]}} + \frac{\sum_{i=1}^{n_C} D_C(i)}{n_C} \\ &\forall \{I=\text{Images}, MM=\text{Multimedia}, C=\text{CSS}\} \\ n &\in R \end{aligned} \quad (3)$$

Actual damage ($D_{m_{actual}}$, defined in Equation 4) is identical to $D_{m_{potential}}$ except it is computed using only the missing embedded resource set R_r .

$$\begin{aligned} D_{m_{actual}} &= \frac{\sum_{i=1}^{n_{I,MM}} D_{[I,MM]}(i)}{n_{[I,MM]}} + \frac{\sum_{i=1}^{n_C} D_C(i)}{n_C} \\ &\forall \{I=\text{Images}, MM=\text{Multimedia}, C=\text{CSS}\} \\ n &\in R_r \end{aligned} \quad (4)$$

In M_m , all embedded resources are treated as equal; all embedded resources are assigned a value of 1 with $weight=1.0$ applied. The potential damage is therefore the number of embedded resources, and the actual damage is the number of missing embedded resources. M_m is the unweighted ratio of missing embedded resources to total embedded resources.

We assign additional weights to important embedded resources at the expense of less important mementos. When a weight w is given to an embedded resource, all n embedded resources lose $\frac{w}{n}$ importance, which redistributes the importance between embedded

resources while keeping the sum of all importance constant. Images receive weights for image size and centrality (Equation 5). We use the pixel area (width \times height) of the image as specified in the HTML and the page size along with a weight for horizontal and vertical central dividing line overlap by the image.

$$\begin{aligned} D_{[I,MM]} &= 1 + \frac{\text{width} * \text{height}}{\text{Page Size (pixels)}} \\ &+ (w_{\text{horizontal}} \iff \text{Overlaps horizontal center}) \\ &+ (w_{\text{vertical}} \iff \text{Overlaps vertical center}) \end{aligned} \quad (5)$$

$w_{\text{horizontal}} = 0.25$
 $w_{\text{vertical}} = 0.25$

Embedded multimedia importance (D_{MM}) is calculated identically to image importance D_I , and we represent both in the same equation $D_{[I,MM]}$. Because size and centrality determine multimedia importance, we omit audio and other non-visual multimedia resources. We also classify Flash movies as multimedia.

Equation 6 outlines the damage from missing stylesheets, including a factor for a style threshold w_{style} .

$$\begin{aligned} D_C &= 1 + w_{style} \iff \\ &(>75\% \text{ non-background in left two columns}) \\ &+ w_{tags} \iff \\ &(\text{tags in the DOM without matching CSS}) \end{aligned} \quad (6)$$

$w_{style} = 0.50$
 $w_{tags} = 0.50$

Our intuition is that a missing important stylesheet will shift content to the left of the page rather than center content in the viewport, we divide a PNG snapshot of a memento into vertical thirds and measure the amount of content in each third. Traditional Web design (and particularly design enabled by stylesheets) evenly distributes content across each of the vertical thirds. If a stylesheet is missing and content appears to be shifted to primarily the left two-thirds, we assume the missing stylesheet was important to the distribution of content on the page.

When detecting content in the PNG snapshot, we use remaining CSS files and the HTML to determine the background color of the page. We measure the number of background and non-background colored pixels, with content being the number of non-background colored pixels. The proportion of non-background colored pixels in each vertical third gives us the amount of content in each partition.

The style threshold is determined as follows:

1. Determine background color

1. ~~have~~
~~controlling~~
~~examples~~
~~HTML~~
~~script~~
~~where~~
~~w X H~~
~~not~~
~~value~~
~~w~~
~~2. where~~
~~w X H~~
~~content~~
~~content~~

- Note: was that missing CSS important? probably an interesting blog post.*
2. Render a PNG snapshot of the page
 3. Divide PNG into vertical third partitions
 4. Calculate number of pixels of the non-background color in each third for the viewport only (we used a 1024x768 viewport) and entire page
 5. If $\leq 75\%$ of the non-background colored pixels are in the left two thirds of the viewport, set $w_{style} = 0$ in Equation 6 (CSS file does not receive a weight)
 6. If $>75\%$ of the non-background colored pixels are in the left two thirds of the viewport and left two thirds of the entire page and a stylesheet is missing, $w_{style} = 0.5$ in Equation 6 (CSS file does receive a weight)

For example, we created two mementos of the URI <http://www.pilotonline.com/> on a local server, one as it appears live (with all stylesheets – Figure 3(a)) and the other with its stylesheets removed (Figure 3(b)). The vertical partitions extend from the top of the PNG snapshot to the bottom. The percent of non-background color pixels in the viewports of our mementos are shown in their respective thirds in Figure 2. Notice that the non-background pixels (text, images, etc.) shift left when the CSS is missing. Intuitively, information is not meant to be displayed like the content in Figure 3(b).

When we consider content outside of the viewport (Figures 4(a) and 4(b)), we see the same shift of content to the left when stylesheets are missing. However, the distribution of content in Figure 4(b) is more evenly distributed because the content has shifted down and fills out the middle and right vertical partitions more than in Figure 3(b). This is an indicator that the stylesheets missing in Figures 3(b) and 4(b) were important.

Along with the style threshold, the presence of tags on the page without a matching style suggests that the missing CSS contained the referenced formatting. If such tags exist without a matching style, $w_{tags} = 0.5$ in the Equation 6.


Embedded multimedia, images, and stylesheets do not account for the entirety of a page's importance and usefulness. We assume that text, as defined by the DOM and included on the page, is available regardless of archival success and therefore does not contribute to the damage calculation.

Equations 1 - 6 are used to compute D_m :

1. Load URI-M with PhantomJS
2. Find Potential Damage $D_{m_{potential}}$ (1)
 - (a) Determine CSS importance D_C (6)
 - (b) Determine Multimedia importance D_{MM} (5)
 - (c) Determine Image importance D_I (5)
3. Determine proportion of unsuccessfully dereferenced embedded resources M_m



(a) We calculated that the non-background color is more evenly distributed between the three vertical partitions of the Pilot Online page with its stylesheet included than when it is missing.




(b) We calculated that the non-background color is most prevalent in the left-most vertical partition of the viewport of the Pilot Online page when it is missing its stylesheet.

Fig. 3 Missing stylesheets causes content to shift left. We show the percent of content in the vertical partitions of the viewport.

4. Find Actual Damage $D_{m_{actual}}$ (same as Step 3, but with only those URI-Ms unsuccessfully dereferenced)
5. Determine total damage $D_m = [0,1]$ (1)

With D_m defined, we revisit the examples presented in Section 2. The values for D_m and M_m are listed in Table 7. Note that the D_m ratings are closer to our empirical human assessment of memento quality than the proportion of the embedded resources that are missing.

Not all pages and page construction methods can be evaluated by this algorithm. An edge case not handled by this algorithm is any page constructed with iframes. Our algorithm uses JavaScript to determine the rendered location of embedded multimedia and im-



(a) When considering the entire page, the content of the page is distributed 33% in the left, 26% in the middle, and 41% in the right partitions when the stylesheet is present.

(b) When considering the entire page, the content of the page is distributed 84% in the left, 15% in the middle, and 1% in the right partitions when the stylesheet is missing.

Fig. 4 Missing stylesheets causes content to shift left. We show the percent of content in the vertical partitions of the page.

Figure	D_m	M_m
1(a)	0.09	0.17
1(b)	0.41	0.24
1(c)	0.36	0.29
1(d)	0.59	0.38
1(e)	0.003	0.20

Table 7 D_m vs M_m for the images in Figure 1. Note $M_m > D_m$ in 2 of 5 cases.

ages. When the embedded media is in a page embedded within another page, our algorithm does not provide the accurate rendered location. For this reason, we exclude iframes from our algorithm. We also exclude missing audio-only multimedia since the sound has no visual impact on the page, and sensory importance beyond sight is not considered in this algorithm.

While D_m includes multimedia calculations, multimedia resources are rarely embedded in our mementos (only observed twice in our entire set of 45,341 URI-Ms). We observed that multimedia is often loaded by JavaScript files embedded in the document object model (DOM); this prevents the multimedia files from being archived since archival crawlers (at the time of this experiment) do not execute client-side JavaScript and therefore do not discover the requested files.

or
frames
?

ΔD_m	Splits						Total
	5-0	4-1	3-2	2-3	1-4	0-5	
1.0							0.00
0.9		1					0.01
0.8							0.00
0.7							0.00
0.6			1				0.01
0.5							0.00
0.4	4	1					0.05
0.3	2	2	3				0.07
0.2		2	1	2	2	1	0.08
0.1	4	16	27	15	12	3	0.77
0.0							0.00
Total	0.10	0.22	0.32	0.17	0.14	0.04	1.0

Table 8 The turker evaluations of the m_2 vs m_3 comparisons when using D_m as a damage measurement.

Turker Assessment	D_m	
	Select m_2	Select m_3
m_2	45	32
m_3	8	14

Table 9 Confusion matrix of the turker assessments of the m_2 vs m_3 comparison test against D_m .

7 Damage in the Archives

Having defined an algorithm for measuring D_m , we measured D_m values for each of the 45,341 URI-Ms from Section 5. We used these measurements to assess D_m 's performance relative to turker assessment and perform damage measurements in the Internet Archive.

7.1 Turker Assessment of D_m

We compared D_m to turker assessment and M_m . As shown in Table 8, D_m agrees with turker assessment of damage 32% of the time, an increase of 18% over M_m . Additionally, 49% tie with a 3-2 or 2-3 split and only 16% of the turker evaluations disagreed with the D_m measure. Turkers agree more consistently when ΔD_m is larger. If we only consider $\Delta D_m > 0.30$, the turkers agree with D_m 71% of the time. However with $\Delta M_m > 0.30$, the turkers agree only 20% of the time.

From the confusion matrix in Table 9, we determine that the accuracy of D_m when comparing m_2 vs m_3 is 0.60, and $F_1 = 0.69$. This is an improvement of 0.14 over the accuracy of M_m and an improvement over the harmonic mean of M_m by 0.14, showing that D_m measures damage closer to turker perception. We also calculated the AUC in a ROC curve for D_m and compared it to M_m and the optimal performance of the m_0 vs m_1 test. As shown in Table 10, D_m has an AUC of 0.584, an increase in 0.108 over M_m , showing that D_m outperforms M_m and is closer to the optimal performance of m_0 vs m_1 (AUC=0.789).

Damage Calculation	AUC	F_1	Accuracy
M_m	0.472	0.55	0.46
D_m	0.584	0.69	0.60
M_{m_0}	0.789	0.88	0.80

Table 10 D_m provides a closer estimate of turker perception of damage and our optimal performance of m_0 vs m_1 than M_m .

7.2 Measuring the Internet Archive

With D_m validated as aligning closer to turker evaluations than M_m , we used D_m to evaluate the Internet Archive's performance. Our measurement shows that only 46% of the 45,341 URI-Ms listed in the 1,861 TimeMaps are complete – that is, 54% of all URI-Ms listed in the Internet Archive TimeMaps we studied are missing at least one embedded resource². In Figure 5, we show the average number of missing embedded resources M_m along with the average calculated damage D_m per URI-M per year.




Fig. 5 The average embedded resources missed per memento per year as compared to damage per memento per year ($D_m=0.128$, $M_m=0.132$).

Because the number of missed mementos is important to M_m and D_m , we investigated the occurrence of missing and successfully dereferenced embedded resources. Most mementos are missing very few embedded resources with most missing 1-10 embedded resources (Figure 6), ($\mu = 1.7$, $\sigma = 4.6$). We calculate that 61% of mementos are missing 3 or fewer embedded resources, and 85% of mementos are missing 6 or fewer embedded resources. While the number of successfully dereferenced embedded resources in mementos is more evenly

² The Internet Archive performs URI canonicalization very well, and is assumed to not be a source of missing resources.

distributed (Figure 7), most mementos have very few embedded resources ($\mu = 17.6$, $\sigma = 86$).




Fig. 6 The distribution of the number of missing embedded resources per URI-M.




Fig. 7 The number of successfully dereferenced resources is more evenly distributed than those missing (Figure 6).

In aggregate, we observed that 45,009 of 292,192 embedded resources were missing, meaning 15% of the embedded resources in the dataset are missing. Of these, 25,848 (57% of the missing URI-Ms) were important, meaning they were assigned an additional weight by D_m (Equations 5 and 6). The average damage of all measured mementos was 0.132.

The yearly D_m goes from an average of 0.16 in 1998 to 0.13 in 2013. That means the Internet Archive is doing a better job (over time) reducing the total mem-

D_m already says we're

onto damage in its collection. However, the number of missing *important* resources (resources with an importance >1 due to added weights) is increasing, going from an average of 1.30 important resources per memento in 1997 to 2.38 important resources per memento in 2013 for an average of 2.05 missing per memento. Meanwhile, the number of unimportant missing embedded resources (damage rating weight ≤ 1) per memento is increasing at a lesser rate, going from 1.35 in 1997 to 1.64 in 2013. This suggests that while the Internet Archive is getting better overall at mitigating damage as much as possible, the archive is missing an increasing number of embedded resources deemed important.

The distribution of file types missing per memento (Figure 8) shows that most URI-Ms are missing ≥ 1 embedded resource and that stylesheets and JavaScript files are missing at higher rates over time. Missing JavaScript may lead to additional missing files (such as multimedia). Images are missing at varying rates per memento over time.




Fig. 8 The number of missed embedded resources per memento per year and MIME type.

7.3 Measuring WebCite

In an effort to measure a less prominent and different type of archive, we used the damage algorithm to determine M_m and D_m of WebCite³[9]. WebCite is different from the Internet Archive's Heritrix crawler in that it is a page-at-a-time archiving tool that creates mementos upon user request. While the Internet Archive's goal is to archive the Web using the Heritrix crawler to identify crawl targets, add their URI-Rs to a crawl frontier, and

crawl the frontier, WebCite and other page-at-a-time archivers allow users to submit URI-Rs for archiving, and WebCite immediately archives the resource⁴. When using a page-at-a-time archival service, the resulting memento contains embedded resources with the same archival datetime [1]. This section identifies our damage measurement of this page-at-a-time archiver and outlines the differences between Heritrix and WebCite.

Our WebCite dataset has 992 mementos in the Time-Maps of our collection of 1,861 URI-Rs. The earliest available memento is from 2007, and the latest is from 2014. Only six mementos are available from 2014; therefore, we will focus on 2007-2013 as the target years of investigation due to the limited number of 2014 mementos, as well as to match the period of observation of the Internet Archive. The average D_m of the collection over all years is $\mu = 0.397$ ($\sigma = 0.194$), and the average M_m is $\mu = 0.176$ ($\sigma = 0.0926$). All of the mementos in this collection are missing at least one embedded resource – 100% of the mementos are incomplete.

As shown in Figure 9, the average D_m in WebCite is increasing over time, going from 0.285 in 2007 to 0.442 in 2013. Meanwhile, the average M_m remains steady, going from 0.135 in 2007 to 0.139 in 2013. Only slight variation occurs, peaking at 0.287 in 2011.

Compared to the Internet Archive, WebCite has a higher damage value as well as is missing a larger percentage of embedded resources. Additionally, the damage rating per memento is higher, indicating that a larger percentage of missing embedded resources are important (3,514 or 41.7%) in WebCite than in the Internet Archive.

WebCite is missing on average 10.1 embedded resources per memento ($\sigma = 8.0$). Across the entire collection, 8,420 of 54,824, or 15.4% of the embedded resources were missing in our investigation. We calculate that 56% of mementos are missing 3 or fewer embedded resources, and 74% of mementos are missing 6 or fewer embedded resources (Figures 10 and 11).

The distribution of file types missing per memento (Figure 12) shows that most URI-Ms are missing ≥ 1 embedded image and CSS resources, on average. WebCite has a lower occurrence of missing stylesheets, but a higher occurrence of missing images. This will impact future work with D_m . If we change the weighting of the importance of missing embedded resources – if we weight missing CSS as having a higher impact on overall D_m , WebCite's collection might have a lower average D_m . However, more investigation is needed before this conclusion may be reached.

⁴ The Internet Archive has recently added an *on-demand* archiving utility at <http://archive.org/web/> under the heading "Save Page Now" [25].

³ <http://webcitation.org/>

(e.g., crawls a single URI-R and not an entire site).




Fig. 9 The average embedded resources missed per memento per year as compared to damage per memento per year ($D_m=0.397$, $M_m=0.176$).




Fig. 10 The distribution of the number of missing embedded resources per URI-M in WebCite.

Our previous investigation [7] showed that WebCite has difficulties when encountering JavaScript and embedded iframes. However, its archiving policies provide immediate results as opposed to crawlers that may incur delay between the time a URI-R is added to the frontier and a memento is created. WebCite's difficulties with JavaScript may contribute to the missing embedded resources if they were loaded through JavaScript.

8 Impact of JavaScript on Damage

As a preliminary investigation of the impact of JavaScript on archival tools, we set up an experiment to use Heritrix and PhantomJS [23] to crawl the same set

↓
already
ref'd




Fig. 11 In WebCite, the number of successfully dereferenced resources is more evenly distributed than those missing (Figure 10).




Fig. 12 The number of missed embedded resources per WebCite memento per year and MIME type.

of URI-Rs and measure the damage difference between the two set of mementos. Our goal is to understand how D_m is impacted by JavaScript by comparing mementos archived by a crawler that can execute JavaScript (PhantomJS) and a crawler that does not execute JavaScript (Heritrix).

8.1 PhantomJS vs Heritrix

Web crawlers operate by starting with a finite set of seed URI-Rs in a frontier – or list of crawl targets – and add to the frontier by extracting embedded URIs in the representations of the URI-R. This allows archival crawlers to discover embedded resources as well as new URI-Rs to crawl while creating mementos.

already
said
this
elsewhere
?

Representations of Web resources are increasingly reliant on JavaScript and other client-side technologies to load embedded resources and control the activity on the client or request additional data or resources (e.g., via Ajax) after the initial page load. Crawlers are unable to discover the resources requested via Ajax and are missing embedded resources which ultimately causes the mementos of the crawled resources to be incomplete and have higher D_m .

To mitigate the impact web developers' practice of using JavaScript and Ajax to load embedded resources, crawlers like Heritrix have constructed approaches for extracting links from embedded JavaScript to be added to crawl frontiers [15] (most recently, Google [6]).

Because archival crawlers' abilities differ from the abilities of browsers, the archives currently hold a representation of the Web from the point of view of crawlers and not Web users. That is, what we archive is increasingly different than what users experience.

8.2 Crawling Deferred Representations

We sampled 50 URI-Rs by randomly generating Bitly.com URIs and identifying the URI-Rs to which the bitly URIs redirected. We then classified the 50 URI-Rs as having deferred representations and crawled the set of URIs with Heritrix and PhantomJS.

During the Heritrix crawl, we used the 50 URI-Rs as a set of seed URIs and allowed Heritrix to create their mementos. The final frontier size of this crawl was 1,588 URIs of embedded resources used to create the mementos. Using our damage algorithm, we measured the damage of the mementos created by Heritrix and found that $\overline{D_m} = 0.148$. Recall that the measured average damage of the Internet Archive was 0.13.

To ensure the crawler executes JavaScript and captures JavaScript-dependent resources during the creation of mementos, we then crawled the 50 URI-Rs with PhantomJS. We recorded the embedded resources needed to create the representation, including those originating from JavaScript. This created a frontier of 3,364 URIs which we used as a seed URI list in Heritrix. We then used Heritrix to create the mementos using only the seed URI list, effectively creating mementos using the frontier list of PhantomJS. The damage of the mementos from this crawl was, on average, 0.1291.

PhantomJS provided a 13.5% improvement to the collection damage over Heritrix. This provides further evidence that JavaScript-dependent representations reduce the quality of mementos due to traditional crawlers' inability to execute JavaScript.

Not only does using PhantomJS provide a larger crawl frontier, but the damage rating of the resulting

mementos is lower. In short, this initial investigation suggests that using PhantomJS mitigates the impact of JavaScript on resources with deferred representations and results in higher-quality mementos.

9 Measuring Archive.today

Archive.today [3] is another page-at-a-time archival service like WebCite. While WebCite does not properly archive deferred representations, Archive.today creates mementos that limit leakage and missing embedded resources typically occurring in other archival services that ignore JavaScript. To study the impact of Archive.today's handling of JavaScript on memento quality, we submitted each of our 1,861 URI-Rs to Archive.today for archiving (much like the way Mink performs user-initiated archiving [17]) to create mementos of each resource.

When Archive.today creates a memento, it modifies the DOM to remove references to embedded resources that were not available at archive time (i.e., embedded resources that returned a non-200 HTTP response code) [22]. This results in a memento that – if created properly – has no missing embedded resources. Additionally, Archive.today obfuscates the URIs of embedded mementos, preventing a reliable mapping from URI-R to URI-M. Due to these two archival practices, the damage algorithm used in this paper is ineffective for determining memento quality. For this reason, we alter the method of measuring the effectiveness of Archive.today's archival process.

We initiated the archiving of each URI-R in our collection by Archive.today. We counted the number of embedded resources that were successfully loaded into the live resources (i.e., returned an HTTP 200 response when their URIs were dereferenced) and compared this number to the number of embedded resources successfully archived by Archive.today. It is worth noting that the delta between the number of embedded resources in live resources and mementos is a measure of neither M_m nor D_m , but is instead a mechanism for understanding memento fidelity.

We found that Archive.today has a delta of $\mu = 19.9$ ($\sigma = 39.2$), meaning on average, Archive.today did not archive 19.9 embedded resources from the live page due to either its inability to archive the resources, or because Archive.today may have deemed the embedded resources not suitable for archiving⁵. A histogram of all deltas is provided in Figure 13.

⁵ Archive.today lists the resources it saves and does not save in its FAQ page at <http://archive.today/faq.html>




Fig. 13 Histogram of the memento vs live resource deltas in Archive.today.

We submitted each URI-R in the collection to WebCite and recorded deltas for the WebCite mementos in the exact way we measured deltas for Archive.today. In this way, we can compare the two page-at-a-time archiver to determine which service creates higher fidelity mementos. WebCite has a delta of $\mu = 21.6$ ($\sigma = 41.7$), which is higher than the average delta of Archive.today. The histogram of the WebCite deltas is provided in Figure 14. The higher average WebCite delta indicates that Archive.today creates higher fidelity mementos than WebCite, likely due to its superior support of JavaScript dependent representations.

10 Conclusions

In this paper, we demonstrated that Web users (as represented by Mechanical Turk Workers) can correctly identify original mementos (m_0 vs m_1) 81% of the time when presented with an original and manually damaged pair of mementos. After randomly selecting 100 URI-Ms from the Internet Archive TimeMaps of 1,861 URI-Rs, we show that turkers' assessment of damage does not match that of M_m ; in fact, their perception of damage more closely aligns with a random selection than with M_m .

To provide a damage metric closer to the perception of Web users, we proposed D_m , a damage calculation al-




Fig. 14 Histogram of the memento vs live resource deltas in WebCite.

gorithm that estimates embedded resource importance to determine the perceived damage of mementos. Using turker evaluations, we showed that D_m aligns with turker perception 32% of the time when considering all ΔD_m values – an improvement of 17% over M_m . If we limit $\Delta D_m > 0.30$, we achieve an agreement of 71%, an improvement of 51% over M_m . We show that the performance of D_m is closer to that of the m_0 vs m_1 test than both M_m and a random selection.

We used D_m to measure the performance of the Internet Archive by measuring $\overline{D_m}$ of 1,861 URI-Rs. The average damage of the Internet Archive collection is 0.13 per memento and is missing 15% of its embedded resources. Mementos are missing 2.05 important resources on average. The Internet Archive has gotten better at mitigating damage over time, reducing D_m from 0.16 (1998) to 0.13 (2013).

Page-at-a-time archivers perform differently than the Internet Archive. We measured mementos of our collection in WebCite, finding that the average damage of the collection is 0.397 per memento and is missing 18% of its embedded resources. Mementos are missing 10.1 resources on average. Even though damage in the Internet Archive is improving, the damage in WebCite is getting worse, increasing D_m from 0.375 (2007) to 0.475 (2013).

We also demonstrate that JavaScript-dependent representations have a detrimental impact on D_m and M_m . By using a crawl strategy in which JavaScript is exe-

cuted during the crawl, damage in the resulting mementos can be improved by 13.5%.

With D_m , archival services can evaluate their performance and the quality of their mementos. The archives could measure a selection of mementos (either randomly sampled or by identifying those missing a proportion of embedded resources, such as $\Delta D_m > 0.30$) for damage to determine whether or not they have been satisfactorily archived. That is, with this algorithm, the archives can provide the greatest damage improvement through targeted repair efforts (e.g., which mementos require additional attention to ensure proper archiving?). Archives can also use historical damage ratings of a URL-R to identify memento improvements or changes.

We also measured the damage of mementos in WebCite, and demonstrated that the damage in the Internet Archive ($\overline{D_m} = 0.128$) is less than that in WebCite ($\overline{D_m} = 0.397$). We know from previous works that WebCite does not archive JavaScript-dependent representations easily. We also measured Archive.today to determine the fidelity of an archival service that makes an effort to use headless browsing to capture JavaScript dependent representations. We found that Archive.today had a delta of 19.9 embedded resources between mementos and live resources, while WebCite had a delta of 21.6. This shows that Archive.today provides a higher fidelity memento than WebCite.

This is a preliminary investigation of memento damage. We have shown that percentage of embedded resources missing is not an accurate representation of damage and have proposed a more accurate metric. Our future work will continue to improve upon the metric by using larger datasets, more turkers, and machine learning to further hone D_m . This will include a refinement of the relative weights of the embedded resources (e.g., the relative importance of CSS vs. images). We will also investigate the cumulative damage rating over time. For example, a logo that never changes over a 5 year period could have increased importance due to its use over multiple mementos. We plan to also measure the damage improvement of mementos if embedded resources are retroactively captured and included in past mementos. This cumulative damage improvement can help identify embedded resources that should be targeted by archives.

11 Acknowledgments

This work supported in part by the National Science Foundation (NSF) (IIS 1009392), the Library of Congress, and the National Endowment for the Humanities (NEH) Digital Humanities Implementation Grant (DHIG) (HK-50181-14).

References

1. Ainsworth, S.G., Nelson, M.L.: Evaluating sliding and sticky target policies by measuring temporal drift in acyclic walks through a web archive. *International Journal on Digital Libraries* pp. 1–16 (2014). DOI 10.1007/s00799-014-0120-4
2. Alnoamany, Y., Alsum, A., Weigle, M., Nelson, M.: Who and What Links to the Internet Archive. In: *Proceedings of the Third International Conference on Theory and Practice of Digital Libraries*, pp. 346–357. ACM (2013). DOI 10.1007/978-3-642-40501-3_35
3. Archive.today: Archive.today. <http://archive.today/> (2013). <http://archive.today/>
4. Banos, V., K. Yunhyong, S.R., Manolopoulos, Y.: CLEAR: a credible method to evaluate website archivability. In: *Proceedings of the 9th International Conference on Preservation of Digital Objects* (2013)
5. Ben Saad, M., Ganarski, S.: Archiving the web using page changes patterns: A case study. In: *Proceedings of the 11th Annual International Joint Conference on Digital Libraries*, pp. 113–122 (2011). DOI 10.1145/1998076.1998098
6. Brunelle, J.F.: Google and JavaScript. <http://ws-d1.blogspot.com/2014/06/2014-06-18-google-and-javascript.html> (2014)
7. Brunelle, J.F., Kelly, M., Weigle, M.C., Nelson, M.L.: The Impact of JavaScript on Archivability. *International Journal of Digital Libraries* TBD(TBD), TBD (TBD). Submitted for publication
8. Denev, D., Mazeika, A., Spaniol, M., Weikum, G.: SHARC: framework for quality-conscious web archiving. *Proceedings of the 35th International Conference on Very Large Data Bases* **2**, 586–597 (2009). DOI 10.1007/s00778-011-0219-9
9. Eysenbach, G., Trudel, M.: Going, going, still there: using the WebCite service to permanently archive cited web pages. *Journal of Medical Internet Research* **7**(5) (2005). DOI 10.2196/jmir.7.5.e60
10. Fawcett, T.: An introduction to ROC analysis. *Pattern recognition letters* **27**(8), 861–874 (2006). DOI 10.1016/j.patrec.2005.10.010
11. Fersini, E., Messina, E., Archetti, F.: Enhancing web page classification through image-block importance analysis. *Information Processing & Management* **44**(4), 1431 – 1447 (2008). DOI 10.1016/j.ipm.2007.11.003
12. GNU: Introduction to GNU Wget. <http://www.gnu.org/software/wget/> (2013)

13. Gray, G., Martin, S.: Choosing a sustainable web archiving method: A comparison of capture quality. *D-Lib Magazine* **19**(5) (2013). DOI 10.1045/may2013-gray
14. Howell, B.A.: Proving Web History: How to Use the Internet Archive. *Journal of Internet Law* **9**(8), 3–9 (2006)
15. Jack, P.: Extractorhtml extract-javascript. <https://webarchive.jira.com/wiki/display/Heritrix/ExtractorHTML+extract-javascript> (2014)
16. Kelly, M., Brunelle, J.F., Weigle, M.C., Nelson, M.L.: On the Change in Archivability of Websites Over Time. In: Proceedings of the Third international conference on Theory and Practice of Digital Libraries, pp. 35–47 (2013). DOI 10.1007/978-3-642-40501-3_5
17. Kelly, M., Nelson, M.L., Weigle, M.C.: Mink: Integrating the live and archived web viewing experience using web browsers and memento. In: Proceedings of ACM/IEEE Digital Libraries (DL). London (2014)
18. Kohlschütter, C., Fankhauser, P., Nejdl, W.: Boilerplate Detection Using Shallow Text Features. In: Proceedings of the Third ACM International Conference on Web Search and Data Mining, pp. 441–450 (2010). DOI 10.1145/1718487.1718542
19. Marshall, C.C., Shipman, F.M.: On the Institutional Archiving of Social Media. In: Proceedings of the 12th ACM/IEEE-CS Joint Conference on Digital Libraries, pp. 1–10 (2012). DOI 10.1145/2232817.2232819
20. Mohr, G., Kimpton, M., Stack, M., Ranitovic, I.: Introduction to Heritrix, an archival quality web crawler. In: Proceedings of the 4th International Web Archiving Workshop (2004)
21. Negulescu, K.C.: Web Archiving @ the Internet Archive. Presentation at the 2010 Digital Preservation Partners Meeting, 2010 <http://www.digitalpreservation.gov/meetings/documents/ndipp10/NDIIPP072110FinalIA.ppt>
22. Nelson, M.L.: 2014-07-14: “Refresh” For Zombies, Time Jumps. <http://ws-dl.blogspot.com/2014/07/2014-07-14-refresh-for-zombies-time.html> (2014)
23. PhantomJS: PhantomJS. <http://phantomjs.org/> (2013)
24. Rademacher, P., Lengyel, J., Cutrell, E., Whited, T.: Measuring the Perception of Visual Realism in Images. In: Rendering Techniques 2001, Eurographics, p. 235247. Springer (2001). DOI 10.1007/978-3-7091-6242-2_22
25. Rossi, A.: Fixing Broken Links on the Internet. <https://blog.archive.org/2013/10/25/fixing-broken-links/> (2013)
26. SalahEldeen, H.M., Nelson, M.L.: Losing my revolution: how many resources shared on social media have been lost? In: Proceedings of the Second international conference on Theory and Practice of Digital Libraries, pp. 125–137 (2012). DOI 10.1007/978-3-642-33290-6_14
27. SalahEldeen, H.M., Nelson, M.L.: Resurrecting My Revolution: Using Social Link Neighborhood in Bringing Context to the Disappearing Web. In: Proceedings of the Third international conference on Theory and Practice of Digital Libraries, pp. 333–345 (2013). DOI 10.1007/978-3-642-40501-3_34
28. Sigursson, K.: Incremental crawling with Heritrix. In: Proceedings of the 5th International Web Archiving Workshop (2005)
29. Singh, R., Bhattacharai, B.D.: Information-theoretic identification of content pages for analyzing user information needs and actions on the multimedia web. In: Proceedings of the 2009 ACM Symposium on Applied Computing, pp. 1806–1810 (2009). DOI 10.1145/1529282.1529686
30. Song, R., Liu, H., Wen, J.R., Ma, W.Y.: Learning block importance models for web pages. In: Proceedings of the 13th international conference on World Wide Web, pp. 203–211 (2004). DOI 10.1145/988672.988700
31. Spaniol, M., Denev, D., Mazeika, A., Weikum, G., Senellart, P.: Data quality in web archiving. In: Proceedings of the 3rd Workshop on Information Credibility on the Web, pp. 19–26. ACM (2009)
32. Spaniol, M., Mazeika, A., Denev, D., Weikum, G.: Catch me if you can: Visual Analysis of Coherence Defects in Web Archiving. In: Proceedings of The 9th International Web Archiving Workshop, pp. 27–37 (2009)
33. Tofel, B.: ‘Wayback’ for Accessing Web Archives. In: Proceedings of the 7th International Web Archiving Workshop (2007)
34. Van de Sompel, H., Nelson, M.L., Sanderson, R., Balakireva, L.L., Ainsworth, S., Shankar, H.: Memento: Time Travel for the Web. Tech. Rep. arXiv:0911.1112, Los Alamos National Laboratory (2009)
35. Yi, L., Liu, B., Li, X.: Eliminating noisy information in web pages for data mining. In: Proceedings of the Ninth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, pp. 296–305 (2003). DOI 10.1145/956750.956785

36. Zhang, X., Lin, W., Xue, P.: Just-noticeable difference estimation with pixels in images. *Journal of Visual Communication and Image Representation* **19**(1), 30–41 (2008). DOI 10.1109/TMM.2013.2268053